

Notes on Parallelizing Matrix-Matrix Multiplication

Robert van de Geijn, Martin Schatz, Tze Meng Low

April 16, 2013

1 A simple implementation

In this assignment, you start with a very simple parallel matrix-matrix multiplication that implements $C = AB + C$ as a sequence of rank-k updates:

$$C = (a_0 \mid a_1 \mid \dots \mid a_{k-1}) \begin{pmatrix} \hat{b}_0^T \\ \hat{b}_1^T \\ \vdots \\ \hat{b}_{k-1}^T \end{pmatrix} = ((\dots((C + a_0 \hat{b}_0^T) + a_1 \hat{b}_1^T) + \dots) + a_{k-1} \hat{b}_{k-1}^T)$$

1.1 Distribution

Consider Figure 1–3, which illustrate how matrices C , A , and B are distributed to a $r \times c$ mesh of (MPI) processes, when $r = c = 3$. They are distributed using an elemental cyclic distribution.

1.2 Routine ParallelRank1

Since the matrix multiplication can be accomplished as a loop around rank-1 updates, I created a routine `ParallelRank1` that implements a rank-1 update with the p th column A multiplying the p th row of B , updating C .

Notice that this routine communicates the p th column of A within rows of nodes and the p th row of B within columns of nodes to achieve the duplication of data illustrated in Figure 2. The duplicated data are in the work arrays `work_A` and `work_B`, respectively, and the local computation is then performed by the call to the BLAS routine `dger_`, which performs a rank-1 update. Notice that `dger_` uses a Fortran interface, meaning that data must be passed by address, which explains some of the ugliness...

1.3 Routine ParallelMMult

The routine `ParallelMMult` in file `ParallelMMult_1.c` simply implements the parallel matrix-matrix multiplication as a loop around calls to `ParallelRank1`.

1.4 The driver routine

The driver routine sets up the MPI environment, creates random matrices, copies (via the utility routine `CopyMatrixGlobalToLocal`) the local part of those random matrices into the local arrays used for the parallel matrix-matrix multiplication, etc. It uses utility routines to create a random matrix and to compare the contents of two matrices.

1.5 Executing

Copy `fvdg/class/CS378S13/PMMult` into a directory of your own. Executing `make run` will compile and test (with a 2×3 mesh of processes). The reported `diff` should be small.

$\beta_{4,0} \ \beta_{4,3} \ \beta_{4,6} \ \cdots$	$\beta_{4,1} \ \beta_{4,4} \ \beta_{4,7} \ \cdots$	$\beta_{4,2} \ \beta_{4,5} \ \beta_{4,8} \ \cdots$
$\alpha_{0,4} \quad \gamma_{0,0} \ \gamma_{0,3} \ \gamma_{0,6} \ \cdots$	$\alpha_{0,4} \quad \gamma_{0,1} \ \gamma_{0,4} \ \gamma_{0,7} \ \cdots$	$\alpha_{0,4} \quad \gamma_{0,2} \ \gamma_{0,5} \ \gamma_{0,8} \ \cdots$
$\alpha_{3,4} \quad \gamma_{3,0} \ \gamma_{3,3} \ \gamma_{3,6} \ \cdots$	$\alpha_{3,4} \quad \gamma_{3,1} \ \gamma_{3,4} \ \gamma_{3,7} \ \cdots$	$\alpha_{3,4} \quad \gamma_{3,2} \ \gamma_{3,5} \ \gamma_{3,8} \ \cdots$
$\alpha_{6,4} \quad \gamma_{6,0} \ \gamma_{6,3} \ \gamma_{6,6} \ \cdots$	$\alpha_{6,4} \quad \gamma_{6,1} \ \gamma_{6,4} \ \gamma_{6,7} \ \cdots$	$\alpha_{6,4} \quad \gamma_{6,2} \ \gamma_{6,5} \ \gamma_{6,8} \ \cdots$
$\vdots \quad \vdots \quad \vdots \quad \vdots \quad \ddots$	$\vdots \quad \vdots \quad \vdots \quad \vdots \quad \ddots$	$\vdots \quad \vdots \quad \vdots \quad \vdots \quad \ddots$
$\beta_{4,0} \ \beta_{4,3} \ \beta_{4,6} \ \cdots$	$\beta_{4,1} \ \beta_{4,4} \ \beta_{4,7} \ \cdots$	$\beta_{4,2} \ \beta_{4,5} \ \beta_{4,8} \ \cdots$
$\alpha_{1,4} \quad \gamma_{1,0} \ \gamma_{1,3} \ \gamma_{1,6} \ \cdots$	$\alpha_{1,4} \quad \gamma_{1,1} \ \gamma_{1,4} \ \gamma_{1,7} \ \cdots$	$\alpha_{1,4} \quad \gamma_{1,2} \ \gamma_{1,5} \ \gamma_{1,8} \ \cdots$
$\alpha_{4,4} \quad \gamma_{4,0} \ \gamma_{4,3} \ \gamma_{4,6} \ \cdots$	$\alpha_{4,4} \quad \gamma_{4,1} \ \gamma_{4,4} \ \gamma_{4,7} \ \cdots$	$\alpha_{4,4} \quad \gamma_{4,2} \ \gamma_{4,5} \ \gamma_{4,8} \ \cdots$
$\alpha_{7,4} \quad \gamma_{7,0} \ \gamma_{7,3} \ \gamma_{7,6} \ \cdots$	$\alpha_{7,4} \quad \gamma_{7,1} \ \gamma_{7,4} \ \gamma_{7,7} \ \cdots$	$\alpha_{7,4} \quad \gamma_{7,2} \ \gamma_{7,5} \ \gamma_{7,8} \ \cdots$
$\vdots \quad \vdots \quad \vdots \quad \vdots \quad \ddots$	$\vdots \quad \vdots \quad \vdots \quad \vdots \quad \ddots$	$\vdots \quad \vdots \quad \vdots \quad \vdots \quad \ddots$
$\beta_{4,0} \ \beta_{4,3} \ \beta_{4,6} \ \cdots$	$\beta_{4,1} \ \beta_{4,4} \ \beta_{4,7} \ \cdots$	$\beta_{4,2} \ \beta_{4,5} \ \beta_{4,8} \ \cdots$
$\alpha_{2,4} \quad \gamma_{2,0} \ \gamma_{2,3} \ \gamma_{2,6} \ \cdots$	$\alpha_{2,4} \quad \gamma_{2,1} \ \gamma_{2,4} \ \gamma_{2,7} \ \cdots$	$\alpha_{2,4} \quad \gamma_{2,2} \ \gamma_{2,5} \ \gamma_{2,8} \ \cdots$
$\alpha_{5,4} \quad \gamma_{5,0} \ \gamma_{5,3} \ \gamma_{5,6} \ \cdots$	$\alpha_{5,4} \quad \gamma_{5,1} \ \gamma_{5,4} \ \gamma_{5,7} \ \cdots$	$\alpha_{5,4} \quad \gamma_{5,2} \ \gamma_{5,5} \ \gamma_{5,8} \ \cdots$
$\alpha_{8,4} \quad \gamma_{8,0} \ \gamma_{8,3} \ \gamma_{8,6} \ \cdots$	$\alpha_{8,4} \quad \gamma_{8,1} \ \gamma_{8,4} \ \gamma_{8,7} \ \cdots$	$\alpha_{8,4} \quad \gamma_{8,2} \ \gamma_{8,5} \ \gamma_{8,8} \ \cdots$
$\vdots \quad \vdots \quad \vdots \quad \vdots \quad \ddots$	$\vdots \quad \vdots \quad \vdots \quad \vdots \quad \ddots$	$\vdots \quad \vdots \quad \vdots \quad \vdots \quad \ddots$

Figure 1: Distribution of C to a 3×3 mesh of processes. Also shown is how a_p and \hat{b}_p are duplicated in order to perform the update $C = C + a_p \hat{b}_p$ as local rank-1 updates.

1.6 Comment

When we run this on a real machine (next week), you will see that this implementation is (weakly) scalable, but achieves very low performance.

2 Optimization 1

The key is to locally call a matrix-matrix multiply rather than a rank-1 update. For this, observe that

$$C = (A_0 \mid A_1 \mid \cdots \mid A_{K-1}) \begin{pmatrix} B_0 \\ B_1 \\ \vdots \\ B_{K-1} \end{pmatrix} = ((\cdots ((C + A_0 B_0) + A_1 B_1) + \cdots) + A_{K-1} B_{K-1})$$

Now, I will describe how to morph `ParallelRank1` into `ParallelRankK`, which will implement $C = C + A_k B_k$, where A_k has b_k columns and B_k has b_k rows.

2.1 Calling sequence

Please make the calling sequence for this new routine

$\begin{array}{cccc} \alpha_{0,0} & \alpha_{0,3} & \alpha_{0,6} & \dots \\ \alpha_{3,0} & \alpha_{3,3} & \alpha_{3,6} & \dots \\ \alpha_{6,0} & \alpha_{6,3} & \alpha_{6,6} & \dots \\ \vdots & \vdots & \vdots & \ddots \end{array}$	$\begin{array}{cccc} \alpha_{0,1} & \alpha_{0,4} & \alpha_{0,7} & \dots \\ \alpha_{3,1} & \alpha_{3,4} & \alpha_{3,7} & \dots \\ \alpha_{6,1} & \alpha_{6,4} & \alpha_{6,7} & \dots \\ \vdots & \vdots & \vdots & \ddots \end{array}$	$\begin{array}{cccc} \alpha_{0,2} & \alpha_{0,5} & \alpha_{0,8} & \dots \\ \alpha_{3,2} & \alpha_{3,5} & \alpha_{3,8} & \dots \\ \alpha_{6,2} & \alpha_{6,5} & \alpha_{6,8} & \dots \\ \vdots & \vdots & \vdots & \ddots \end{array}$
$\begin{array}{cccc} \alpha_{1,0} & \alpha_{1,3} & \alpha_{1,6} & \dots \\ \alpha_{4,0} & \alpha_{4,3} & \alpha_{4,6} & \dots \\ \alpha_{7,0} & \alpha_{7,3} & \alpha_{7,6} & \dots \\ \vdots & \vdots & \vdots & \ddots \end{array}$	$\begin{array}{cccc} \alpha_{1,1} & \alpha_{1,4} & \alpha_{1,7} & \dots \\ \alpha_{4,1} & \alpha_{4,4} & \alpha_{4,7} & \dots \\ \alpha_{7,1} & \alpha_{7,4} & \alpha_{7,7} & \dots \\ \vdots & \vdots & \vdots & \ddots \end{array}$	$\begin{array}{cccc} \alpha_{1,2} & \alpha_{1,5} & \alpha_{1,8} & \dots \\ \alpha_{4,2} & \alpha_{4,5} & \alpha_{4,8} & \dots \\ \alpha_{7,2} & \alpha_{7,5} & \alpha_{7,8} & \dots \\ \vdots & \vdots & \vdots & \ddots \end{array}$
$\begin{array}{cccc} \alpha_{2,0} & \alpha_{2,3} & \alpha_{2,6} & \dots \\ \alpha_{5,0} & \alpha_{5,3} & \alpha_{5,6} & \dots \\ \alpha_{8,0} & \alpha_{8,3} & \alpha_{8,6} & \dots \\ \vdots & \vdots & \vdots & \ddots \end{array}$	$\begin{array}{cccc} \alpha_{2,1} & \alpha_{2,4} & \alpha_{2,7} & \dots \\ \alpha_{5,1} & \alpha_{5,4} & \alpha_{5,7} & \dots \\ \alpha_{8,1} & \alpha_{8,4} & \alpha_{8,7} & \dots \\ \vdots & \vdots & \vdots & \ddots \end{array}$	$\begin{array}{cccc} \alpha_{2,2} & \alpha_{2,5} & \alpha_{2,8} & \dots \\ \alpha_{5,2} & \alpha_{5,5} & \alpha_{5,8} & \dots \\ \alpha_{8,2} & \alpha_{8,5} & \alpha_{8,8} & \dots \\ \vdots & \vdots & \vdots & \ddots \end{array}$

Figure 2: Distribution of A to a 3×3 mesh of processes.

```

void ParallelRankK(
    // global matrix dimensions m, n, k
    int global_m, int global_n, int global_k,
    // the global index of the current columns of A and current rows of B
    // to be used for the rank-b_k update
    int p,
    // the block size for this rank-k update
    int bk,
    // address of where local A, B, and C are stored, as well as their
    // local "leading dimensions"
    double *local_A, int local_ldimA,
    double *local_B, int local_ldimB,
    double *local_C, int local_ldimC,
    // communicator for the row in which this node is a member
    MPI_Comm comm_row,
    // communicator for the column in which this node is a member
    MPI_Comm comm_col )

```

(Notice the new parameter bk.)

Store this first implementation in file ParallelRankK.2.c.

$\beta_{0,0} \ \beta_{0,3} \ \beta_{0,6} \ \cdots$ $\beta_{3,0} \ \beta_{3,3} \ \beta_{3,6} \ \cdots$ $\beta_{6,0} \ \beta_{6,3} \ \beta_{6,6} \ \cdots$ $\vdots \quad \vdots \quad \vdots \quad \ddots$	$\beta_{0,1} \ \beta_{0,4} \ \beta_{0,7} \ \cdots$ $\beta_{3,1} \ \beta_{3,4} \ \beta_{3,7} \ \cdots$ $\beta_{6,1} \ \beta_{6,4} \ \beta_{6,7} \ \cdots$ $\vdots \quad \vdots \quad \vdots \quad \ddots$	$\beta_{0,2} \ \beta_{0,5} \ \beta_{0,8} \ \cdots$ $\beta_{3,2} \ \beta_{3,5} \ \beta_{3,8} \ \cdots$ $\beta_{6,2} \ \beta_{6,5} \ \beta_{6,8} \ \cdots$ $\vdots \quad \vdots \quad \vdots \quad \ddots$
$\beta_{1,0} \ \beta_{1,3} \ \beta_{1,6} \ \cdots$ $\beta_{4,0} \ \beta_{4,3} \ \beta_{4,6} \ \cdots$ $\beta_{7,0} \ \beta_{7,3} \ \beta_{7,6} \ \cdots$ $\vdots \quad \vdots \quad \vdots \quad \ddots$	$\beta_{1,1} \ \beta_{1,4} \ \beta_{1,7} \ \cdots$ $\beta_{4,1} \ \beta_{4,4} \ \beta_{4,7} \ \cdots$ $\beta_{7,1} \ \beta_{7,4} \ \beta_{7,7} \ \cdots$ $\vdots \quad \vdots \quad \vdots \quad \ddots$	$\beta_{1,2} \ \beta_{1,5} \ \beta_{1,8} \ \cdots$ $\beta_{4,2} \ \beta_{4,5} \ \beta_{4,8} \ \cdots$ $\beta_{7,2} \ \beta_{7,5} \ \beta_{7,8} \ \cdots$ $\vdots \quad \vdots \quad \vdots \quad \ddots$
$\beta_{2,0} \ \beta_{2,3} \ \beta_{2,6} \ \cdots$ $\beta_{5,0} \ \beta_{5,3} \ \beta_{5,6} \ \cdots$ $\beta_{8,0} \ \beta_{8,3} \ \beta_{8,6} \ \cdots$ $\vdots \quad \vdots \quad \vdots \quad \ddots$	$\beta_{2,1} \ \beta_{2,4} \ \beta_{2,7} \ \cdots$ $\beta_{5,1} \ \beta_{5,4} \ \beta_{5,7} \ \cdots$ $\beta_{8,1} \ \beta_{8,4} \ \beta_{8,7} \ \cdots$ $\vdots \quad \vdots \quad \vdots \quad \ddots$	$\beta_{2,2} \ \beta_{2,5} \ \beta_{2,8} \ \cdots$ $\beta_{5,2} \ \beta_{5,5} \ \beta_{5,8} \ \cdots$ $\beta_{8,2} \ \beta_{8,5} \ \beta_{8,8} \ \cdots$ $\vdots \quad \vdots \quad \vdots \quad \ddots$

Figure 3: Distribution of B to a 3×3 mesh of processes.

2.2 From ParallelRank1 to ParallelRankK_2

Work arrays First, you will want to make it so that `work.A` can hold b_k columns and `work.B` can hold b_k columns. Hint: take rows of B but store them as columns of `work.B`!!! This will make your life easier.

Communication In your first implementation, perform separate broadcasts to communicate each of the b_k columns of A and b_k rows of B by putting a loop around the broadcasts in `ParallelRank1`.

Local computation Also, put a loop around the rank-1 updates (calls to `dger_`).

2.3 From ParallelMMult_1 to ParallelMMult_2

Copy `ParallelMMult_1.c` into `ParallelMMult_2.c` (but don't change the name or calling sequence). Change the routine to call `ParallelRankK`. You will have to change the loop to increment in steps of `bk`. How do you deal with the fact that in the last iteration you may not have a full set of `bk` columns in A (or rows in B)?

Change the `Makefile` and see if you get the right answer.

3 Optimization 2

3.1 From ParallelRankK_2 to ParallelRankK_3

Copy from one file to a new file.

Local computation Replace the loop with calls to `dger_` with a single call to the matrix-matrix multiplication routine `dgemm_`. (An example of how to call this is in the driver. Remember: `work_B` hold the transpose of the rows of B with which you are computing!)

Change the `Makefile` and see if you get the right answer.

4 Optimization 3

4.1 From ParallelRankK_3 to ParallelRankK_4

Copy from one file to a new file.

Communication Replace the loop with calls to `MPI_Bcast` with a single call to an appropriate collective communication. What else do you need to do?

Change the `Makefile` and see if you get the right answer.